

# EECS 368/468 – Lab 1

## Programming Massively Parallel Processors with CUDA

---

### *Matrix Multiplication in CUDA*

In this first assignment, you will work in **teams of two** to write a matrix multiplication code in CUDA. The lab is due on **Friday Jan 23 at 11:59pm**. The late submission policies are explained in the syllabus. Please follow the instructions below to do the lab.

#### **Install the CUDA SDK at your Wilkinson Lab account**

1. Login to a machine at Wilkinson Lab (Tech M338). If you wish to work remotely (e.g., via ssh), the Wilkinson Lab hostnames are:

alfred.eecs.northwestern.edu  
bane.eecs.northwestern.edu  
batgirl.eecs.northwestern.edu  
batman.eecs.northwestern.edu  
clayface.eecs.northwestern.edu  
cobblepott.eecs.northwestern.edu  
freeze.eecs.northwestern.edu  
gordon.eecs.northwestern.edu  
gotham.eecs.northwestern.edu  
harley.eecs.northwestern.edu  
huntress.eecs.northwestern.edu  
hush.eecs.northwestern.edu  
joker.eecs.northwestern.edu  
killercroc.eecs.northwestern.edu  
madhatter.eecs.northwestern.edu  
nightwing.eecs.northwestern.edu  
poisonivy.eecs.northwestern.edu  
ras.eecs.northwestern.edu  
riddler.eecs.northwestern.edu  
robin.eecs.northwestern.edu  
scarecrow.eecs.northwestern.edu  
selina.eecs.northwestern.edu  
twoface.eecs.northwestern.edu

2. Setup the CUDA environment. If you use csh or tcsh:  
`source /usr/local/cuda-5.0/cuda-env.csh`  
or if you use the bash shell:  
`./usr/local/cuda-5.0/cuda-env.sh`
3. Install SDK in your home directory  
`cp -r /home/hardav/NVIDIA_CUDA-5.0_Samples ~/;`  
`cd ~/NVIDIA_CUDA-5.0_Samples;`  
`make`
4. Test by finding out the characteristics of your GPU device  
`cd ~/NVIDIA_CUDA-5.0_Samples/1_Uutilities/deviceQuery;`  
`./deviceQuery`

### Install the labs scaffold at your Wilkinson Lab account

5. Download the labs-scaffold tarball from blackboard into some directory in your Wilkinson Lab account.
6. Untar the code scaffold tarball:  

```
tar xvf EECS468-CUDA-Labs-Scaffold.tgz
```

### Install lab1 at your Wilkinson Lab account

7. Download the lab1 tarball from blackboard into some directory in your Wilkinson Lab account.
8. Install the lab1 tarball and compile the lab sources:  

```
cd EECS468-CUDA-Labs  
tar xvf EECS468-CUDA-Lab1.tgz  
make
```

### Complete the assignment

9. **Remember:** every time you login to the Wilkinson lab to program in CUDA you need to setup the CUDA environment. If you use csh or tcsh:  

```
source /usr/local/cuda-5.0/cuda-env.csh
```

  
or if you use the bash shell:  

```
./usr/local/cuda-5.0/cuda-env.sh
```

#### 10. **Assignment Task:**

Edit the following two functions

`MatrixMulOnDevice(...)` in the source file `matrixmul.cu`

`MatrixMulKernel(...)` in the source file `matrixmul_kernel.cu`

to implement matrix multiplication on the GPU device. **Do not change any other piece of the source code.** The size of the matrix is defined such that one thread block will be sufficient to compute the entire solution matrix.

The source code you will be working on is at :

`EECS468-CUDA-Labs/labs/src/lab1`

Compiling your code produces the executable:

`EECS468-CUDA-Labs/labs/bin/linux/release/lab1`

11. There are several modes of operation for the application.
  - a. No arguments: the application will create two randomly initialized matrices to multiply. After the device multiplication is invoked, it will compute the correct solution matrix using the CPU, and compare that solution with the device-computed solution. If it matches (within a certain tolerance), it will print out "Test PASSED" to the screen before exiting.
  - b. One argument: the application will use the random initialization to create the input matrices, and write the device-computed output to the file specified by the argument.

- c. Two arguments: the application will initialize the two input matrices with the values found in the files provided as arguments. No output is written to file.
- d. Three arguments: the application will read its inputs from the files provided by the first two arguments, and write its output to the file provided in the third.

Note that if you wish to use the output of one run of the application as an input, you must delete the first line in the output file, which displays the accuracy of the values within the file. The value is not relevant for this application.

12. **Hand in your solution:**

- a. Create a tarball of your solution, which includes the contents of the `lab1` source folder provided, with all the changes and additions you made to the source files. Please make sure you do a `make clobber` before submitting; don't include object code or executables, as these are typically large files:

```
cd EECS468-CUDA-Labs/labs/src/lab1;
make clobber
tar cvfz EECS468-lab1-YourNames.tgz *
```

where you should replace `YourNames` with the **concatenated last names of the team members**.

- b. Submit your solution tarball via blackboard. Only one student in a team needs to submit. If you submit multiple times, only the latest submission will be graded. Please do all submissions from the same account.

Good luck!